

Graphene: Efficient Interactive Set Reconciliation Applied to Blockchain Propagation

19210240055


袁和昕



Background




- P2p distributed systems
- End-points need to talk to each other




Background: Txns



- Transactions (txns) are transfers of money
- Unvalidated txns are broadcast




Background: Blocks




- Transactions (txns) are transfers of money
- Unvalidated txns are broadcast

Background: Blocks



- Transactions (txns) are transfers of money
- Unvalidated txns are broadcast
- Blocks are comprised of txns



Background: Mempool





- Blocks are comprised of txns
- Each peer in the network has a pool of unvalidated txns, called the **mempool**
- Peers clear out txns from their mempool

Motivation

- Many distributed systems require synchronization of records among processes
- Blockchains are just the latest example
 - Replicas in distributed databases
 - Distributed sensors
 - Security certifications
- Must solve **set reconciliation**



Setup




- **Goal:** Send as little data as possible over the wire


Problem Definition



- Given a block of txns from Alice, and a set of txns at Bob, determine:



- The subset of Bob's txns that are in the block



- The subset of txns that Bob is missing

Contributions



- **A new protocol** that solves which elements in a set M stored by a receiver are members of a subset $N \subseteq M$
- **Extension of our protocol** where some of the elements of N are missing
- **Efficient search algorithm for parameterizing an IBLT**
- Evaluation using **open-source deployment** in the real-world, mathematical analysis, and **simulation**


Bloom Filters

- Bloom filters represent a set of n items
- binary array T (initial value is 0), $\frac{-n \log_2 f}{\ln^2 2}$ bits
- k hash functions h_1, \dots, h_k
 - $\forall x \in S, 1 \leq i \leq k, T[h_i(x)] = 1, \forall x \in S, 1 \leq i \leq k$ 1
 - $T[h_i(x)] == 1, 1 \leq i \leq k$ 2

Bloom Filters

- The False Positive Rate (FPR) is tunable
 - More bits will lower the FPR
- Number of FPs we will observe approximately follow a binomial distribution with two parameters:
 - n : number of items to test for membership
 - p : probability of failure

If **FPR** = $\frac{1}{m - n}$, then we expect
1 transaction from mempool to
falsely appear to be in the $m - n$ txs



Invertible Bloom Lookup Tables



- IBLTs are a generalization of Bloom Filters
 - Instead of a bit, cells include a count and actual content
- IBLTs I
 - j items, $c = j\tau$ rows, $k + 1$ hash functions, k sub-tables of size c/k

row	count	keySum	value	checkSum
0	1	0x6170706c65	0b0011	100f7a
1	1	0x6170706c65	0b0011	100f7a
...


Invertible Bloom Lookup Tables



- We can separate the key-value pairs through the table rows where the **count** is 1 and recalculate the insertion position to delete the key-value pairs from the table.
By gradually deleting the table row where the **count** is 1, the original set S can be recovered from the IBLT table I .


row	count	keySum	Value	checkSum
0	1	0x6170706c65	0b0011	100f7a
1	1	0x6170706c65	0b0011	100f7a
...

Invertible Bloom Lookup Tables



- IBLTs support **subtraction**
 - IBLTs must be **the same size** for subtraction
 - Subtraction recovers **symmetric difference**, $(S_A - S_B) \cup (S_B - S_A)$
- If subtraction recovers the entire symmetric difference, then we say that the **subtraction decoded**

Invertible Bloom Lookup Tables




row	count	keySum	Value	checkSum
0	1	0x6170706c65	0b0011	100f7a
1	-1	0x4300754343	044326	100e1b
...

Graphene




- There is something to optimize!
- The figure shows that we don't need a low FPR for Bloom filter
- The IBLT can help us recover from the mistakes made by the Bloom filter




Graphene



- A occurs in X with probability at least β .
- Given a block of txns from Alice (blue)



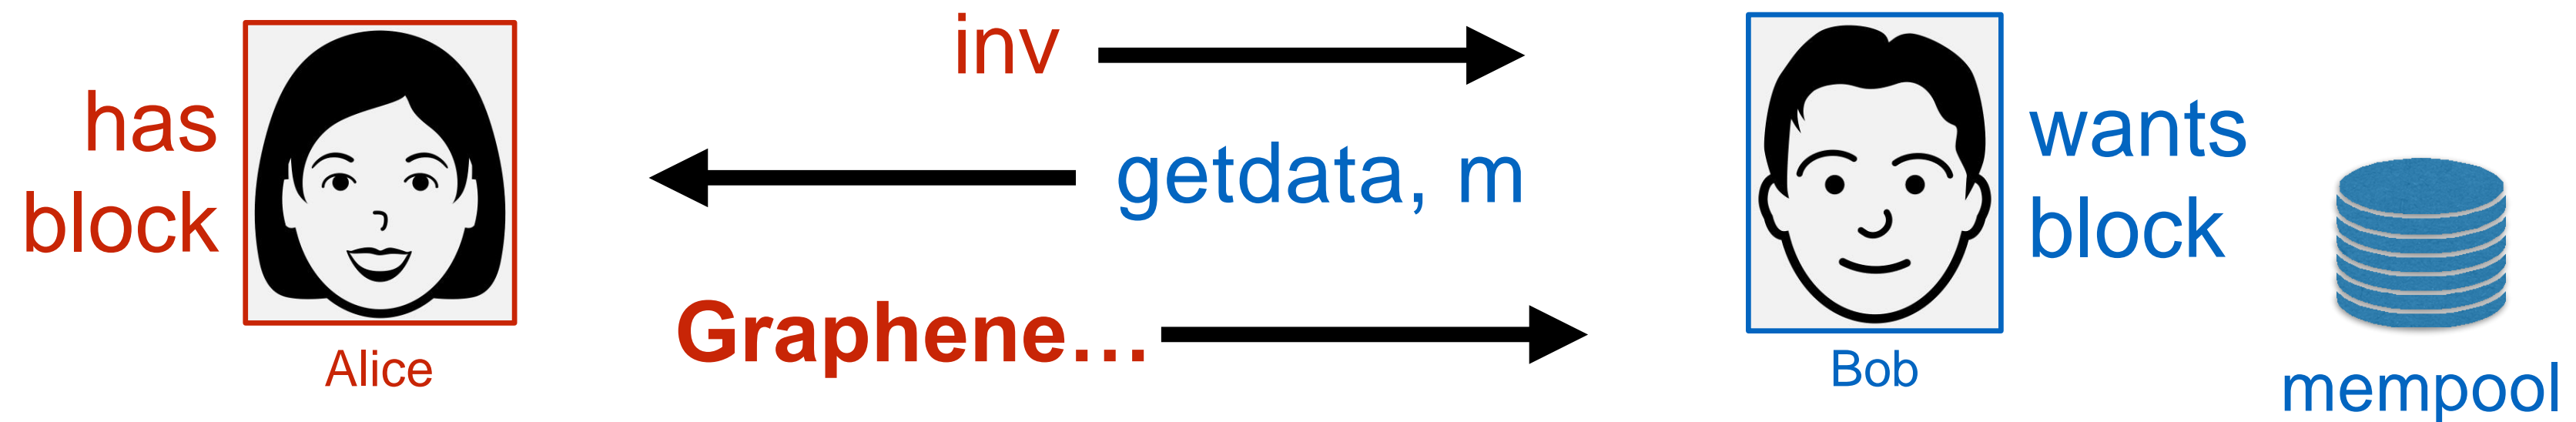
- Protocol 1: The subset of Bob's txns that are in the block (purple)



- Protocol 2: The subset of txns that Bob is missing (red)

Compact Blocks

- inv: block headers
- getdata: requests the block if needed
- Alice's block: n txs
Bob's mempool: m txs



Graphene: Protocol 1




- The block is a subset of the mempool

Sender



Receiver



block

in block


NOT in block

mempool

Graphene: Protocol 1



- The block is a subset of the mempool




Graphene: Protocol 1



- The block is a subset of the mempool

Sender



$$f_S = \frac{a}{m - n}$$





Figure 4: [Protocol 1] Passing m mempool transactions through S results in a FPs (in dark blue). A green outline illustrates $a^* > a$ with β -assurance, ensuring IBLT I decodes.

Graphene: Protocol 1



- The block is a subset of the mempool



Receiver








Figure 4: [Protocol 1] Passing m mempool transactions through S results in a FPs (in dark blue). A green outline illustrates $a^* > a$ with β -assurance, ensuring IBLT I decodes.

Graphene: Protocol 1






- The block is a subset of the mempool



Graphene: Protocol 1



- The block is a subset of the mempool



Graphene: Protocol 1



- The block is a subset of the mempool

Sender



$$f_S = \frac{a}{m-n}, a^* = (1 + \delta)a$$

$$T_I = r\tau(1 + \delta)a, T_{BF} = \frac{-n \ln f_S}{8(\ln 2)^2}$$

$$T = T_{BF} + T_I = \frac{-n \ln(\frac{a}{m-n})}{8(\ln 2)^2} + r\tau(1 + \delta)a$$

block

Bloom
Filter S

IBLT
I

Graphene: Protocol 1



- The block is a subset of the mempool

Sender



block

Bloom
Filter S

IBLT
 I

$$T(a) = T_{BF} + T_I = \frac{-n \ln(\frac{a}{m-n})}{8(\ln 2)^2} + r\tau(1 + \delta)a$$

$$a = n/(8r\tau \ln^2 2), \delta = 0$$

accurate only for $a > 100$

Graphene: Protocol 1



- **THEOREM 1: Derivation of a^***

Let m be the size of a mempool that contains all n transactions from a block. If a is the number of false positives that result from passing the mempool through Bloom filter S with FPR f_S , then $a^ \geq a$ with probability β when*

$$a^* = (1 + \delta)a,$$

$$\text{where } \delta = \frac{1}{2}(s + \sqrt{s^2 + 8s}) \text{ and } s = \frac{-\ln(1-\beta)}{a}$$

Graphene: Protocol 1



- **THEOREM 1: Derivation of a^* , PROOF**

LEMMA 1: Let A be the sum of i independent Bernoulli trials A_1, \dots, A_i , with mean $\mu = E[A]$. Then for $\delta > 0$

$$Pr[A \geq (1 + \delta)\mu] \leq \text{Exp}\left(-\frac{\delta^2}{2 + \delta}\mu\right)$$

Starting from the well-known Chernoff bound

$$\begin{aligned} Pr[A \geq (1 + \delta)\mu] &\leq \left(\frac{e^\delta}{(1 + \delta)^{1+\delta}}\right)^\mu \\ &= \text{Exp}(\mu(\delta - (1 + \delta)\ln(1 + \delta))) \\ &\leq \text{Exp}\left(\mu\left(\delta - (1 + \delta)\left(\frac{2\delta}{2 + \delta}\right)\right)\right) \\ &= \text{Exp}\left(\frac{-\delta^2}{2 + \delta}\mu\right) \end{aligned}$$

$$\ln(1 + x) \geq \frac{x}{1+x/2} = \frac{2x}{2+x} \text{ for } x > 0$$

Graphene: Protocol 1



- **THEOREM 1: Derivation of a^* , PROOF**

There are $m - n$ potential false positives that pass through S . They are a set A_1, \dots, A_{m-n} of independent Bernoulli trials such that $\Pr[A_i = 1] = f_S$. Let $\sum_{i=1}^{m-n} A_i = A$ and $\mu = E[A] = f_S(m-n) = a$. From Lemma 1, we have

$$\Pr[A \geq (1 + \delta)\mu] \leq \text{Exp}\left(-\frac{\delta^2}{2 + \delta}\mu\right).$$

$$\begin{aligned}\beta &= 1 - \text{Exp}\left(-\frac{\delta^2}{2 + \delta}a\right) \\ \delta &= \frac{1}{2}(s + \sqrt{s^2 + 8s}), \text{ where } s = \frac{-\ln(1-\beta)}{a}\end{aligned}$$

Graphene: Protocol 2





- The block is *not* a subset of the mempool


Sender



Receiver




Δ



Can't
decode!

missing txns!



Graphene: Protocol 2



- The block is *not* a subset of the mempool





Figure 5: [Protocol 2] Passing m transactions through S results in z positives, obscuring a count of x TPs (purple) and y FPs (in dark blue). From z , we derive $x^* < x$ with β -assurance (in green).

Receiver



missing txns!

Bloom
Filter R

true negatives

Graphene: Protocol 2



- The block is *not* a subset of the mempool




Figure 6: [Protocol 2] From our bound $m - x^* > m - x$ with β -assurance (in yellow), we can derive a bound for the false positives from S as $y^* > y$ with β -assurance outlined in green.



Bloom
Filter R

$$f_R = \frac{b}{n - x}$$



Graphene: Protocol 2



- The block is *not* a subset of the mempool

Sender





Figure 5: [Protocol 2] Passing m transactions through S results in z positives, obscuring a count of x TPs (purple) and y FPs (in dark blue). From z , we derive $x^* < x$ with β -assurance (in green).

Graphene: Protocol 2



- The block is *not* a subset of the mempool

Sender





Figure 6: [Protocol 2] From our bound $m - x^* > m - x$ with β -assurance (in yellow), we can derive a bound for the false positives from S as $y^* > y$ with β -assurance outlined in green.

Graphene: Protocol 2






- The block is *not* a subset of the mempool

Sender



Receiver



Graphene: Protocol 2



- The block is *not* a subset of the mempool




Figure 6: [Protocol 2] From our bound $m - x^* > m - x$ with β -assurance (in yellow), we can derive a bound for the false positives from S as $y^* > y$ with β -assurance outlined in green.

Graphene: Protocol 2



- **Parameterizing b**

We show below that $y^* = (1 + \delta)y$. Thus, for protocol 2 the total size is

$$T(b) = \frac{z \ln(\frac{b}{n-x^*})}{8 \ln^2 2} + r\tau(1 + \delta)b$$

$$b = z/(8r\tau \ln^2 2)$$

Graphene: Protocol 2



- Using z to parameterize R and J




Figure 5: [Protocol 2] Passing m transactions through S results in z positives, obscuring a count of x TPs (purple) and y FPs (in dark blue). From z , we derive $x^* < x$ with β -assurance (in green).




Figure 6: [Protocol 2] From our bound $m - x^* > m - x$ with β -assurance (in yellow), we can derive a bound for the false positives from S as $y^* > y$ with β -assurance outlined in green.

Graphene: Protocol 2



- **THEOREM 2:**

Let m be the size of a mempool containing $0 \leq x \leq n$ transactions from a block. Let $z = x + y$ be the count of mempool transactions that pass through S with FPR f_S , with true positive count x and false positive count y . Then $x^ \leq x$ with probability β when*

$$x^* = \arg \min_{x^*} Pr[x \leq x^*; z, m, f_S] \leq 1 - \beta.$$

$$\text{where } Pr[x \leq k; z, m, f_S] \leq \sum_{i=0}^k \left(\frac{e^{\delta_k}}{(1 + \delta_k)^{1+\delta_k}} \right)^{(m-k)f_S}$$

$$\text{and } \delta_k = \frac{z - k}{(m - k)f_S} - 1.$$

Graphene: Protocol 2



- **THEOREM 2 PROOF:**

Let Y_1, \dots, Y_{m-x} be independent Bernoulli trials representing transactions not in the block that might be false positives such that $\Pr[Y_i = 1] = f_s$.

Let $\sum_{i=1}^{m-x} Y_i = Y$ and $y = E[Y]$.

For a given value x , we can compute $\Pr[Y \geq y]$, the probability of at least y false positives passing through the sender's Bloom filter. We apply a Chernoff bound:

$$\Pr[y; z, x, m] = \Pr[Y \geq (1 + \delta)\mu] \leq \left(\frac{e^\delta}{(1 + \delta)^{1+\delta}} \right)^\mu$$

Graphene: Protocol 2



- **THEOREM 2 PROOF:**

where $\delta > 0$, and $\mu = E[Y] = (m - x)f_S$. By setting $(1 + \delta)\mu = z - x$ and solving for δ , we have

$$(1 + \delta)(m - x)f_S = z - x$$

$$\delta = \frac{z - x}{(m - x)f_S} - 1.$$

Graphene: Protocol 2



- **THEOREM 2 PROOF:**

$$y = z - x$$

$$Pr[x \leq k; z, m, f_S] = \sum_{i=0}^k Pr[y; z, k, m]$$

$$\leq \sum_{i=0}^k \left(\frac{e^{\delta_k}}{(1 + \delta_k)^{1+\delta_k}} \right)^{(m-k)f_S}$$

$$\text{where } \delta_k = \frac{z-k}{(m-k)f_S} - 1.$$

$$\arg \min_{x^*} Pr[x \leq x^*; z, m, f_S] \leq 1 - \beta$$

- **THEOREM 3:**

Let m be the size of a mempool containing $0 \leq x \leq n$ transactions from a block. Let $z = x + y$ be the count of mempool transactions that pass through S with FPR f_S , with true positive count x and false positive count y . Then $y^ \geq y$ with probability β when*

$$y^* = (1 + \delta)(m - x^*)f_S,$$

$$\text{where } \delta = \frac{1}{2}(s + \sqrt{s^2 + 8s}) \text{ and } s = \frac{-\ln(1 - \beta)}{(m - x^*)f_S}$$

Graphene: Protocol 2



- **THEOREM 3 PROOF:**

We find $y^ = z - x^* \geq y$ by applying Lemma 1 to $\sum_{i=1}^{m-x^*} Y_i = Y$, the sum of $m - x^*$ independent Bernoulli such that might be false positives such that $\Pr[Y_i = 1] = f_S$ trials and $\mu = (m - x^*)f_S$*

$$\Pr[Y \geq (1 + \delta)\mu] \leq \text{Exp} \left(-\frac{\delta^2}{2 + \delta} \mu \right)$$

$$\beta = 1 - \text{Exp} \left(-\frac{\delta^2}{2 + \delta} (m - x^*) f_S \right)$$

$$\delta = \frac{1}{2}(s + \sqrt{s^2 + 8s}), \text{ where } s = \frac{-\ln(1 - \beta)}{(m - x^*) f_S}.$$

Graphene: Protocol 2



- THEOREM 3 PROOF:

$$y^* = (1 + \delta)(m - x^*)f_S$$





Figure 20: [Simulation, Protocol 2] The fraction of Monte Carlo experiments where $y^* > y$ via Theorem 3 compared to a desired bound of $\beta = 239/240$ (shown as a red dotted line).

ENHANCING IBLT PERFORMANCE



Let $H = (V, X, k)$ be a k -partite, k -uniform hypergraph, composed of a set of c vertices. Let $V = V_1 \cup \dots \cup V_k$, where each V_i is a subset of c/k vertices (we enforce that c is divisible by k). X is a set of j hyperedges, each connecting k vertices, one from each of the V_i .

ENHANCING IBLT PERFORMANCE



		IBLT	
		count	value
Hash 1	1	2	$j1 \otimes j2$
	2	2	$j3 \otimes j4$
	3	1	$j5$
Hash 2	4	2	$j1 \otimes j2$
	5	2	$j3 \otimes j5$
	6	1	$j4$
Hash 3	7	2	$j1 \otimes j2$
	8	2	$j4 \otimes j5$
	9	1	$j3$

Hypergraph equivalent

edge	connected vertices
$j1$	$v1, v4, v7$
$j2$	$v1, v4, v7$
$j3$	$v2, v5, v9$
$j4$	$v2, v6, v8$
$j5$	$v3, v5, v8$

} 2-core

$V = V1 \cup V2 \cup V3$
 $V1 = \{v1, v2, v3\}$
 $V2 = \{v4, v5, v6\}$
 $V3 = \{v7, v8, v9\}$

- j items and hyper-edges
- c cells and vertices
- k hash functions and vertices connecting each edge

$$H_{j,p} = \{(V, X, k) \mid E[\text{decode}((V, X, k))] \geq p, |X| = j\}$$

$$\arg \min_{(V, X, k) \in \mathcal{H}_{j,p}} |V|$$

ALGORITHM 1: IBLT-Param-Search

```
01 SEARCH( $j, k, p$ ):  
02    $c_l = 1$   
03    $c_h = c_{max}$   
04    $trials = 0$   
05    $success = 0$   
06    $L = (1 - p)/5$   
07   WHILE  $c_l \neq c_h$ :  
08      $trials += 1$   
09      $c = (c_l + c_h)/2$   
10     IF  $\text{decode}(j, k, c)$ :  
11        $success += 1$   
12      $\text{conf} = \text{conf\_int}(success, trials)$   
13      $r = success/trials$   
14     IF  $r - \text{conf} \geq p$ :  
15        $c_h = c$   
16     IF  $(r + \text{conf} \leq p)$ :  
17        $c_l = c$   
18     IF  $(r - \text{conf} > p - L)$  and  $(r + \text{conf} < p + L)$ :  
19        $c_l = c$   
20   RETURN  $c_h$ 
```

ENHANCING IBLT PERFORMANCE



- Larger but computation is faster
- Create a universal lookup table

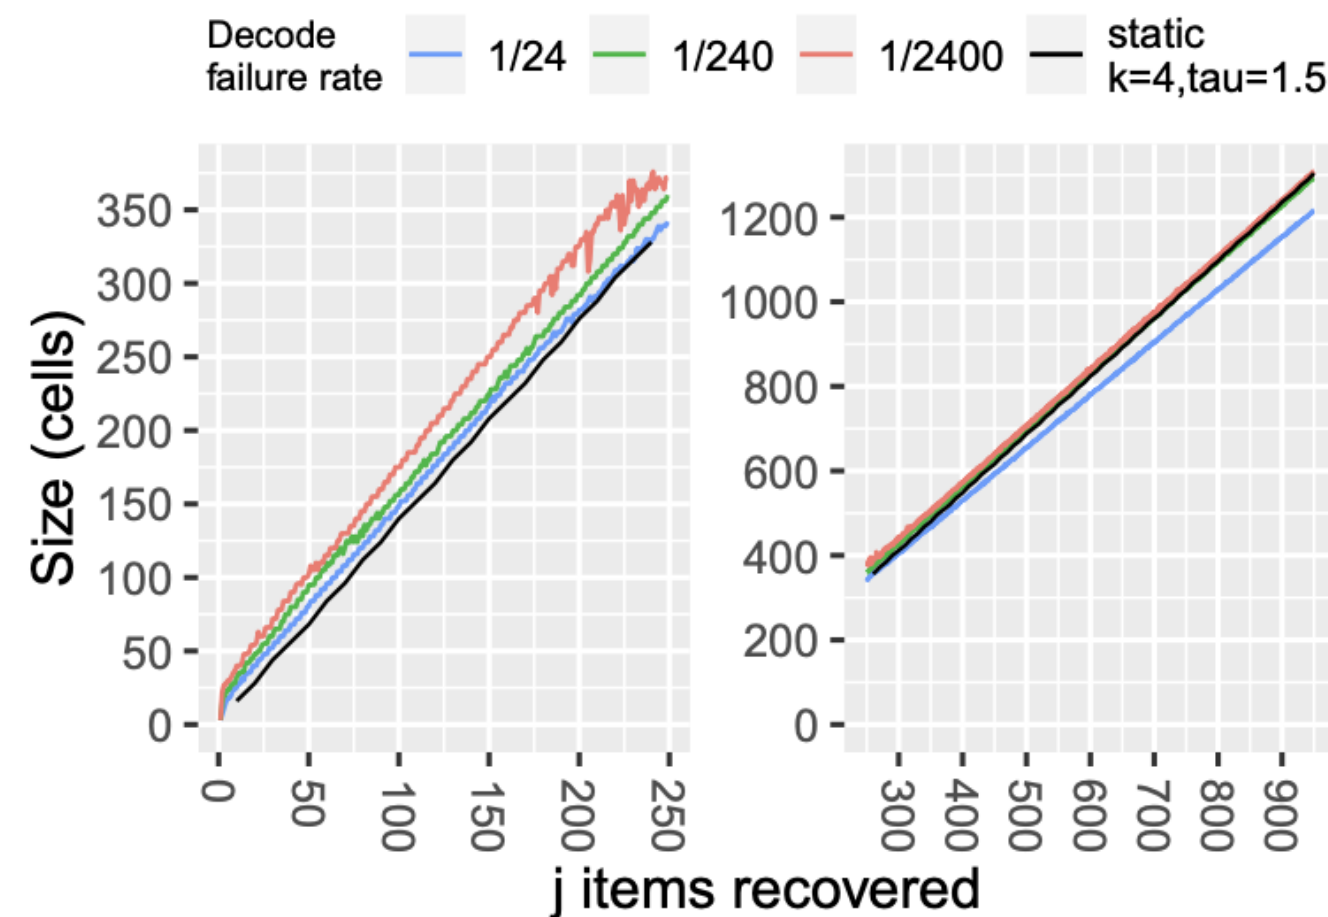


Figure 10: Size of optimal IBLTs (using Alg. 1) given a desired decode rate; with a statically parameterized IBLT ($k = 4$, $\tau = 1.5$) in black. For clarity, the plot is split on the x -axis. Decode rates are shown in Fig. 7.

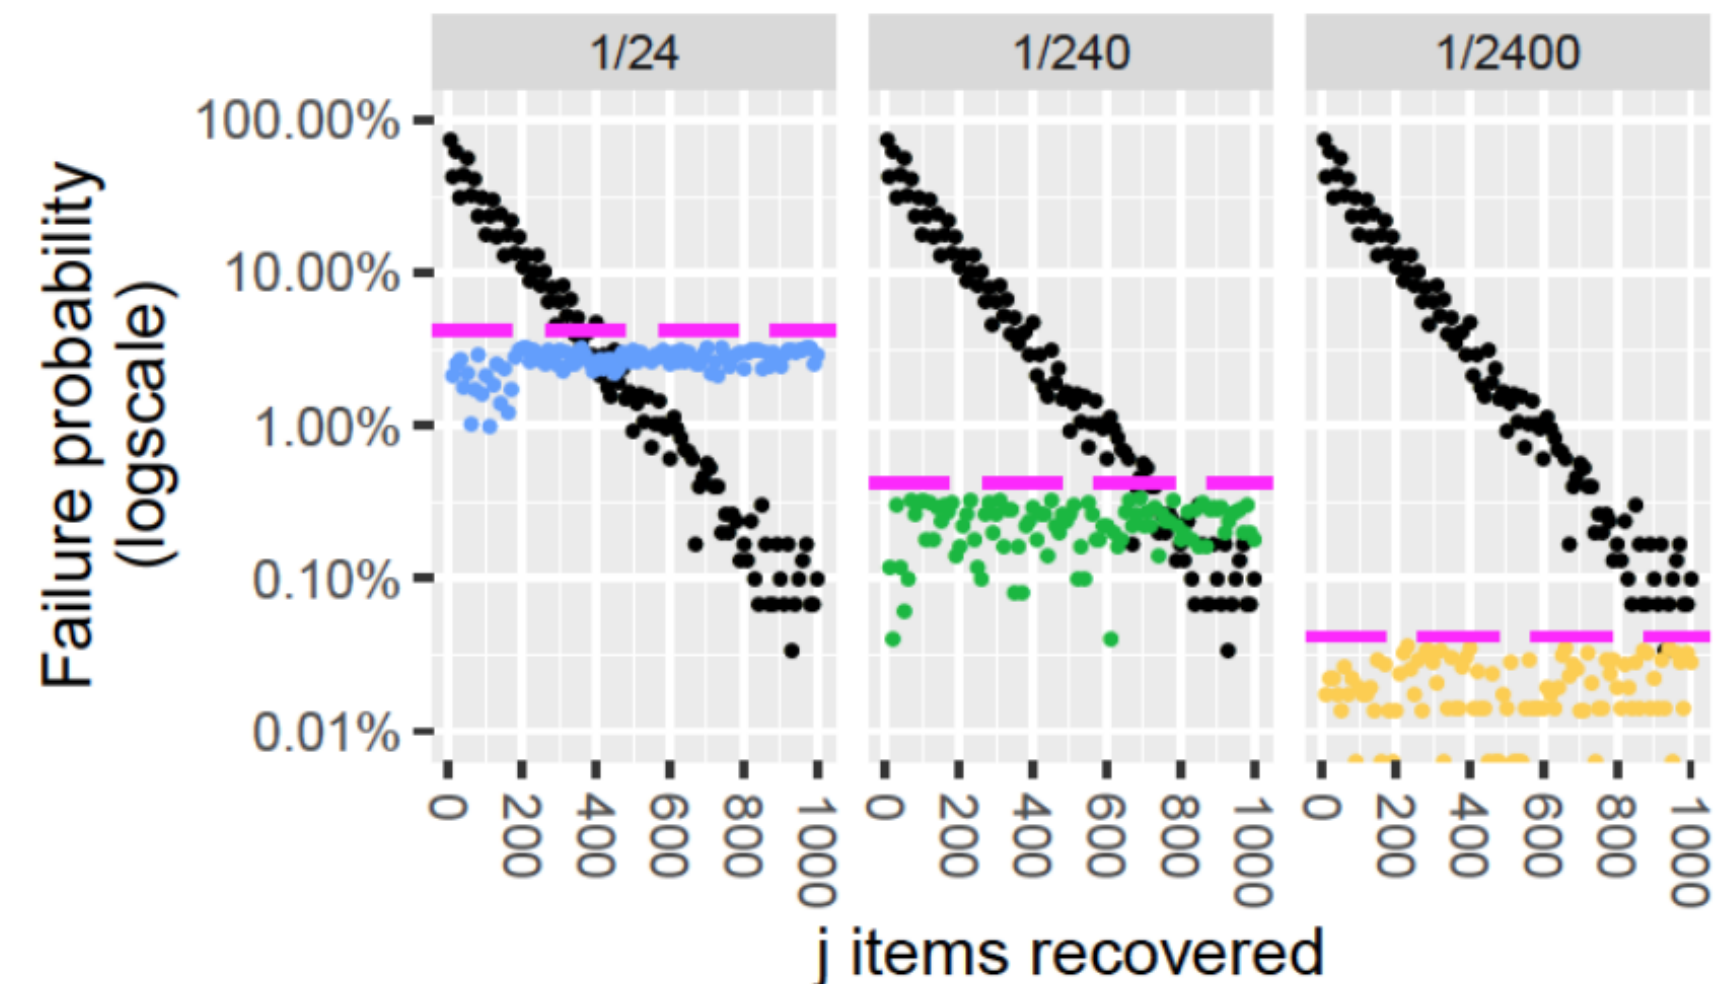


Figure 7: Parameterizing an IBLT statically results in poor decode rates. The black points show the decode failure rate for IBLTs when $k = 4$ and $\tau = 1.5$. The blue, green and yellow points show decode failure rates of optimal IBLTs, which always meet a desired failure rate on each facet (in magenta). Size shown in Fig. 10.

ENHANCING IBLT PERFORMANCE



- Ping-Pong Decoding

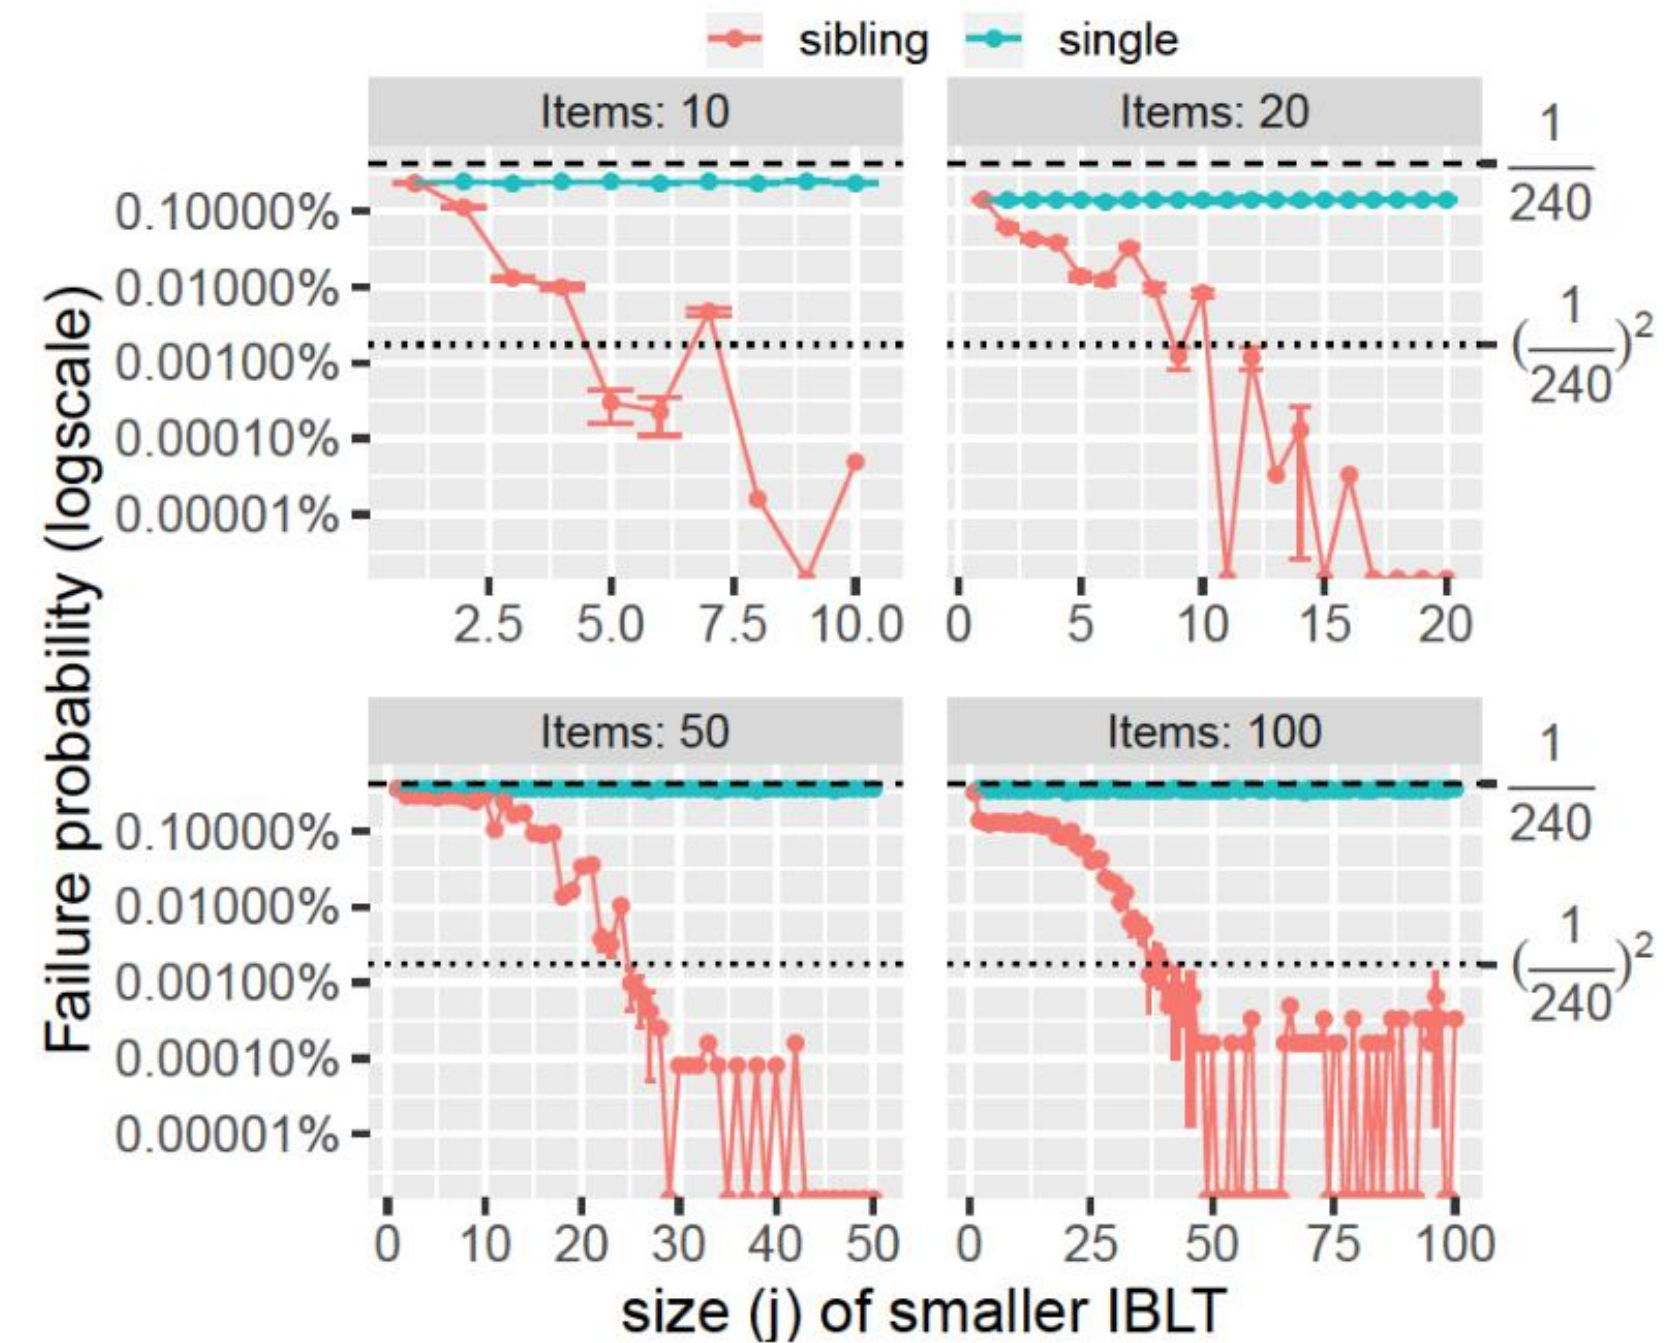



Figure 11: Decode rate of a single IBLT (parameterized for a $1/240$ failure rate) versus the improved *ping-pong decode* rate from using a second, smaller IBLT with the same items.

Open-Source Deployment




- Deployed on Bitcoin Cash network via the Bitcoin Unlimited client
- 1,431 nodes
- Fraction of the size of previous work
- Deploying a protocol requires real engagement with the community
- Adversarial thinking is critical
- Mempools are in-sync less often than expected



Evaluation



- Three block sizes in terms of number of txns
- The receiver's mempool contains
 - All transactions in the block
 - Additional txns as a multiple of the block size
- Improvement with block size



Summary

Thank you!

- **SYSTEMS ISSUES**
 - Security Considerations
 - Transaction Ordering Costs
 - Reducing Processing Time
- **Limitations**
- **CONCLUSION**

